

Turn any document into a **clean table** in one command line.

Six steps. Ten minutes. One API key.



Samy Yusif

AI Products Section Head

Install it. One Python package, nothing else.

```
pip install llm
```

A command line client that talks to OpenAI, Anthropic, Gemini and local models with the same syntax. You learn one interface, not one vendor's interface.



Samy Yusif

AI Products Section Head

Store the API key once, never inside a command.

```
llm keys set openai
```

It prompts you and writes the key to a local file. Nothing after this asks for it again, so no key ever ends up pasted into a command you might share or screenshot.



Samy Yusif

AI Products Section Head

Describe the table you want in plain words.

```
llm --schema 'item, qty int, rate float'
```

That is the entire schema language. A column name, then int, float or bool where you need a type. No JSON Schema to hand write and nothing to validate.



Samy Yusif

AI Products Section Head

Point it at the document and ask for every row.

```
llm --schema-multi 'item, qty int' -a boq.pdf
```

The -a flag attaches the file. --schema-multi returns many rows instead of one object, which is what you want from anything holding a list: a BOQ, a log, an invoice batch.



Samy Yusif

AI Products Section Head

Save it as a template so you never retype it.

```
llm --schema-multi 'item, qty int' --save boq
```

The next document becomes one short line: `llm -t boq -a next.pdf`. This is where the ten minutes pays back. The schema is now an asset you own, not a message in a lost chat.



Samy Yusif

AI Products Section Head

Pull every run back out as one JSON array.

```
llm logs --schema t:boq --data-key items
```

Every response is logged to a local SQLite file. Add `--data-array` to get a single array you can pipe into `sqlite-utils`, or paste straight into a spreadsheet.



Samy Yusif

AI Products Section Head

Reconcile before you trust one extracted number.

Count the rows against the printed lines, then sum them against the printed total. If the difference is not exactly zero, the extraction failed, however tidy the table looks.



Samy Yusif

AI Products Section Head

A scanned PDF is a picture, not a document.

The model reads pixels, so faint stamps, rotated pages and handwriting disappear silently. Confirm your model accepts application/pdf, and run scans through OCR first.



Samy Yusif

AI Products Section Head

Try it first on a document you already know by heart.

You will catch the mistakes immediately. That is the only honest way to learn how far the extraction can be trusted on a document you have never read.



Samy Yusif

AI Products Section Head