

AI that **triages RFIs**, studied the way a product is studied.

Twelve stages, four on architecture. It routes, never answers.



Samy Yusif

AI Products Section Head

RFIs are answered late, and a late RFI becomes a claim.

The register is not the pain. The pain is that everything queues at one priority, so the question sitting on a critical activity waits behind one that could have waited a month.



Samy Yusif

AI Products Section Head

Only the RFIs on the critical path turn into money.

$\text{cost} = \text{days late} \times \text{delay cost per day}$

Apply it to critical activities only. Counting every overdue RFI inflates the case and gets it rejected by the first person who knows the programme. Size the real subset, not the register.



Samy Yusif

AI Products Section Head

Route with a machine. Answer with an engineer.

Classification and prioritisation have a defensible right answer, so a model suits them. Answering an RFI is a contractual position taken on behalf of the company. That does not move.



Samy Yusif

AI Products Section Head

The change lands at intake, not at the answer.

As-is, an RFI arrives and joins one queue. To-be, it arrives tagged by discipline, linked to its drawing and clause, and marked whether it touches a critical activity. The same people answer it.



Samy Yusif

AI Products Section Head

The return is recovered days, not saved clicks.

`value = days recovered x delay cost/day`

Recovered days on critical activities only, minus build and run cost. Time saved on sorting is real but small. If the case only closes on sorting time, the product is not worth building.



Samy Yusif

AI Products Section Head

A classifier, two resolvers, and one sorted queue.

Classify discipline and urgency. Resolve the drawing and clause the RFI refers to. Resolve the schedule activity it blocks. Output an ordered queue. Nothing in this diagram writes an answer.



Samy Yusif

AI Products Section Head

Three systems have to agree on one identifier.

`RFI -> drawing/clause -> activity id`

The RFI register, the drawing and specification set, and the programme. Without a shared code the link is guesswork dressed as automation. That integration is the product, not a prerequisite to it.



Samy Yusif

AI Products Section Head

Two confidence scores, never one combined number.

route score and criticality score

Getting the discipline wrong wastes a day. Missing that an RFI sits on the critical path costs a claim. One blended score buries the expensive error underneath the cheap one.



Samy Yusif

AI Products Section Head

It sorts the queue. It never drafts, closes or sends.

No auto-reply, no auto-close, nothing leaves for the client without a human. Say it in the architecture diagram, not in a footnote, because that line is what makes the whole thing approvable.



Samy Yusif

AI Products Section Head

The dangerous failure is confident wrong routing.

An RFI carrying a claim, sent to the wrong discipline and quietly aged. Worse, criticality read from a programme not updated in six weeks, which makes the flag decorative while looking authoritative.



Samy Yusif

AI Products Section Head

The schedule must be current, or none of this is real.

The criticality link is worth exactly as much as the last update. Add a full audit trail, because RFIs are contractual documents and automated handling of them will eventually be examined.



Samy Yusif

AI Products Section Head

Replay last year's RFIs, where the outcome is already known.

prioritise recall on the claim flag

Missing an RFI that carried a claim costs far more than one false alarm. Set a high recall floor on that flag, accept the false positives, and set a separate accuracy target for routing.



Samy Yusif

AI Products Section Head

Run it in shadow before it is allowed to decide.

The machine routes and the technical office routes, in parallel, neither seeing the other. Compare for a full cycle. The disagreements are the training data and the argument for going live.



Samy Yusif

AI Products Section Head

Adoption here is how often nobody edits the routing.

```
adoption = accepted / total routed
```

If the team overrides the suggestion constantly, the tool is not adopted, it is tolerated. That one ratio tells you more than any usage dashboard a vendor will offer you.



Samy Yusif

AI Products Section Head

Median days to close, on critical items, before and after.

$$\text{value} = (\text{base days} - \text{now}) \times \text{cost/day}$$

Median, not average, because one monstrous RFI will otherwise carry the entire result. And measured on the critical subset you defined in stage two, not on the register as a whole.



Samy Yusif

AI Products Section Head

Notice how much of this is not about the model.

Ten of the twelve stages would still exist if a person did the routing. That is what separates a product decision from a demo, and it is where the money is actually decided.



Samy Yusif

AI Products Section Head