AI that triages RFIs, studied the way a product is studied.

Twelve stages. Formulas only. It routes, it never answers.

Samy Yusif

AI Products Section Head

RFIs are answered late, and a late RFI becomes a claim.

The register is not the pain. The pain is that everything queues at the same priority, so the question sitting on a critical activity waits behind one that could have waited a month.

Samy Yusif
AI Products Section Head

Only the RFIs on the critical path turn into money.

```
cost = days late x delay cost per day
```

Apply it to critical activities only. Counting every overdue RFI inflates the case and gets it rejected by the first person who knows the programme. Size the real subset, not the register.

Samy Yusif
AI Products Section Head

Route with a machine. Answer with an engineer.

Classification and prioritisation have a defensible right answer, so they suit a model. Answering an RFI is a contractual position taken on behalf of the company. That does not move.

Samy Yusif
AI Products Section Head

The change lands at intake, not at the answer.

As-is, an RFI arrives and joins one queue. To-be, it arrives tagged by discipline, linked to its drawing and clause, and marked whether it touches a critical activity. Same people answer it.

Samy Yusif
AI Products Section Head

The return is recovered days, not saved clicks.

```
value = days recovered x delay cost/day
```

Recovered days on critical activities only, minus build and run cost. Time saved on sorting is real but small. If the case only closes on sorting time, the product is not worth building.

Samy Yusif
AI Products Section Head

A classifier, two lookups, and one human queue.

Classify discipline and urgency, resolve the drawing and specification clause referenced, then look up the linked activity in the programme. Output is a sorted queue, nothing more.

Samy Yusif
AI Products Section Head

The dangerous failure is confident wrong routing.

An RFI carrying a claim, sent to the wrong discipline and quietly aged. Worse, criticality read from a programme that has not been updated in six weeks, which makes the flag decorative.

Samy Yusif
AI Products Section Head

The schedule must be current, or none of this is real.

The criticality link is only as honest as the last update. Plus an audit trail, because RFIs are contractual documents and any automated handling of them will eventually be examined.

Samy Yusif
AI Products Section Head

Replay last year's RFIs, where the outcome is known.

`prioritise recall on the claim flag`

Missing an RFI that carried a claim costs far more than one false alarm. So set a high recall floor on that flag, accept the false positives, and set a separate accuracy target for routing.

Samy Yusif
AI Products Section Head

Run it in shadow before it is allowed to decide.

The machine routes and the technical office routes, in parallel, without seeing each other. Compare for a full cycle. Disagreements are the training data and the argument for going live.

Samy Yusif
AI Products Section Head

Adoption here is how often nobody edits the routing.

```
adoption = accepted / total routed
```

If the team overrides the suggestion constantly, the tool is not adopted, it is tolerated. That single ratio tells you more than any usage dashboard the vendor will offer you.

Samy Yusif
AI Products Section Head

Median days to close, on critical items, before and after.

```
value = (base days - now) x cost/day
```

Median, not average, because one monstrous RFI will otherwise carry the whole result. And measured on the critical subset you defined in stage two, not on the register as a whole.

Samy Yusif
AI Products Section Head

Notice how much of this is not about the model.

Ten of the twelve stages would still exist if a person did the routing. That is what separates a product decision from a demo, and it is where the money is actually decided.

Samy Yusif
AI Products Section Head